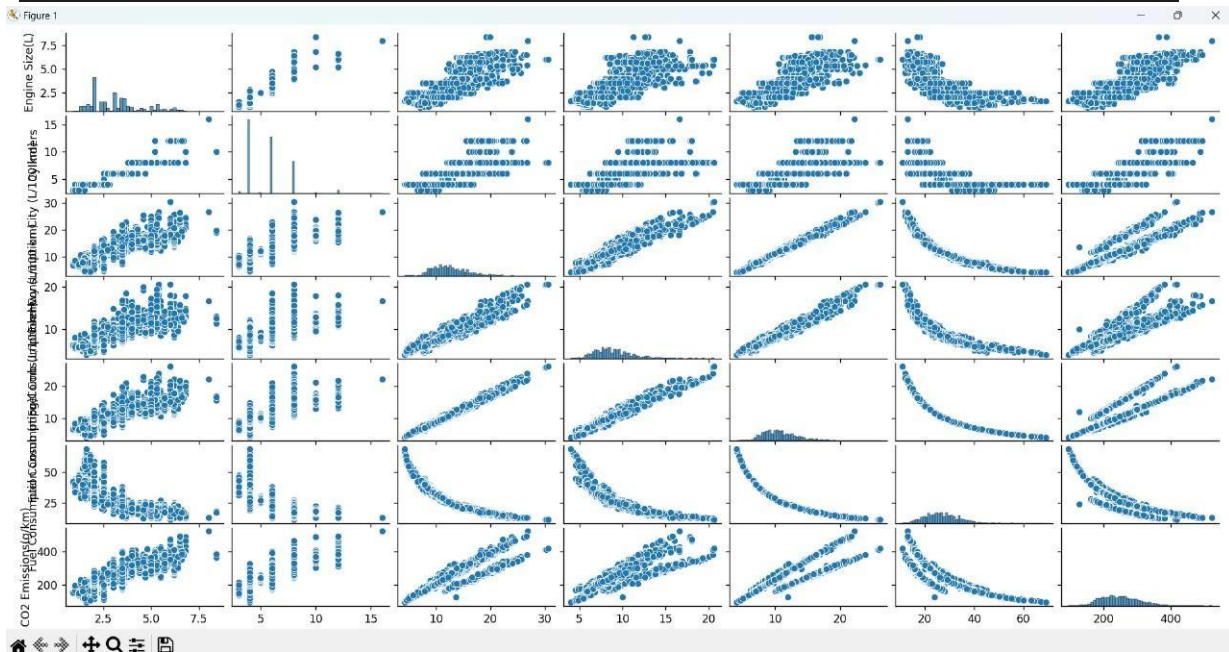


# Report:

Name	ID
Ali Bedair	20206036
Youssef Abdellah	20215043

## Analysis:


```
sns.pairplot(dataset[numeric_features])  
plt.show()
```



```

Matrix_Correlation = dataset[numeric_features].corr()
plt.figure(figsize=(9, 7))
plt.title("This is correlation heatmap of numeric features")
sns.heatmap(Matrix_Correlation, annot=True, fmt='.3f', linewidths=1, cmap='coolwarm')
plt.show()

```



```
print("\nMissing values in each column:")
print(dataset.isnull().sum())

numeric_features = dataset.select_dtypes(include=['float64', 'int64']).columns
summary_info = dataset.describe().to_string()
print("\nSummary info of Dataset analysis:")
print(summary_info)
```

Missing values in each column:

Make	0
Model	0
Vehicle Class	0
Engine Size(L)	0
Cylinders	0
Transmission	0
Fuel Type	0
Fuel Consumption City (L/100 km)	0
Fuel Consumption Hwy (L/100 km)	0
Fuel Consumption Comb (L/100 km)	0
Fuel Consumption Comb (mpg)	0
CO2 Emissions(g/km)	0
Emission Class	0
dtype: int64	

Summary info of Dataset analysis:

	Engine Size(L)	Cylinders	Fuel Consumption City (L/100 km)	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)	Fuel Consumption Comb (mpg)	CO2 Emissions(g/km)
count	7385.000000	7385.000000	7385.000000	7385.000000	7385.000000	7385.000000	7385.000000
mean	3.160068	5.619030	12.556534	9.041706	10.975071	27.481652	250.504699
std	1.354170	1.820307	3.500274	2.224456	2.892506	7.231879	58.512679
min	0.900000	3.000000	4.200000	4.000000	4.100000	11.000000	96.000000
25%	2.000000	4.000000	10.100000	7.500000	8.900000	22.000000	208.000000
50%	3.000000	6.000000	12.100000	8.700000	10.600000	27.000000	246.000000
75%	3.700000	6.000000	14.600000	10.200000	12.600000	32.000000	288.000000
max	8.400000	16.000000	30.600000	20.600000	26.100000	69.000000	522.000000

Selected features:

```
mean = independent_features_train.mean(axis=0)
standard_deviation = independent_features_train.std(axis=0)
Scaled_numeric_train = (independent_features_train - mean) / standard_deviation
Scaled_numeric_test = (independent_features_test - mean) / standard_deviation

print("Scaled Training Features (First 5 rows):\n", Scaled_numeric_train.head())
print("Scaled Test Features (First 5 rows):\n", Scaled_numeric_test.head())
```

Scaled Training Features (First 5 rows):

	Engine Size(L)	Fuel Consumption City (L/100 km)
6214	2.098659	1.684512
900	1.212508	1.079238
77	-0.855178	-0.621294
4488	0.917124	0.502786
6955	0.474048	-0.073665

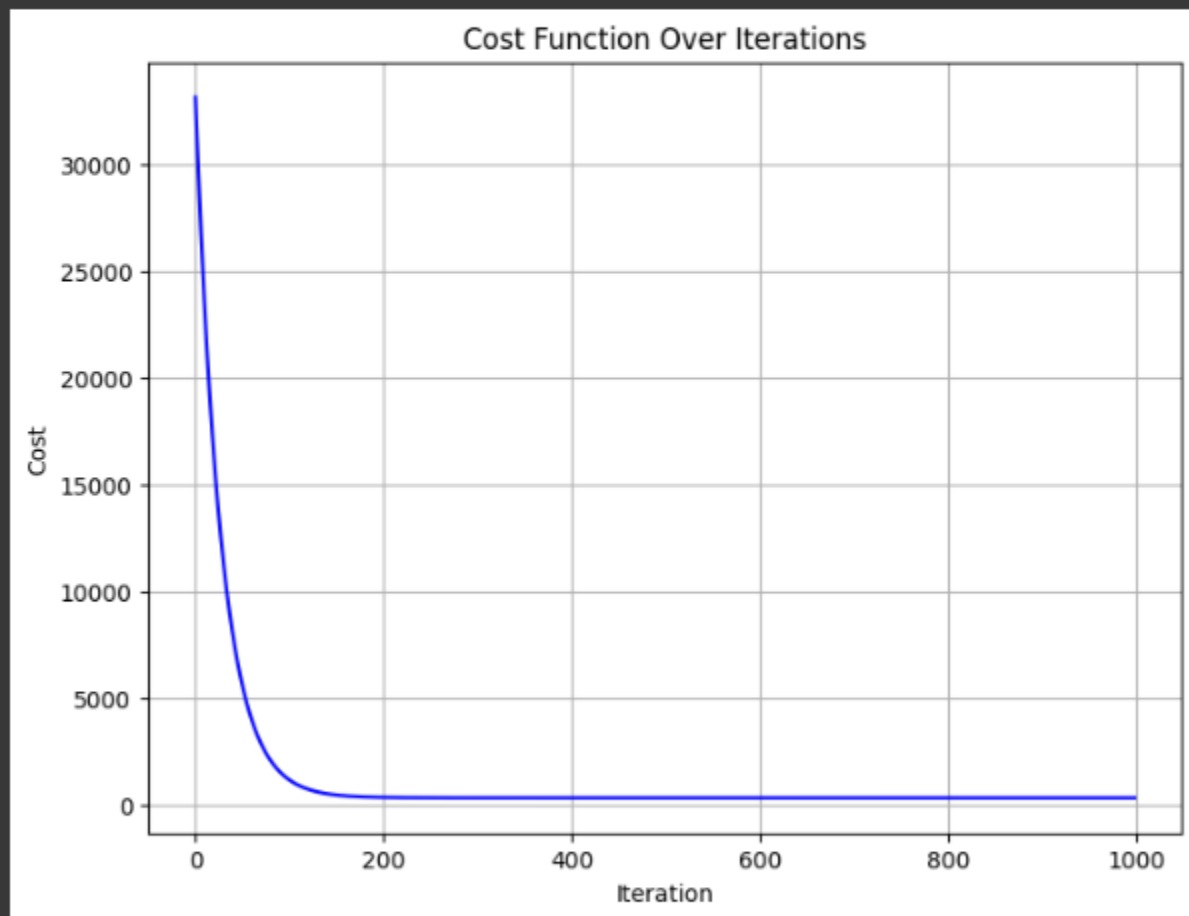
Scaled Test Features (First 5 rows):

	Engine Size(L)	Fuel Consumption City (L/100 km)
5340	-0.559795	-0.390713
3783	-0.559795	-0.967165
1582	0.252511	-0.073665
309	0.326356	0.185738
1849	1.138662	0.099270

Cost function per iterations:

```
# Making predictions on the test set
predictions_test = X_test.dot(theta)
print("\nFirst 6 predictions on the test set:")
print(predictions_test[:6])

# test set using Scikit-learn's R2 score
r2 = r2_score(emission_test_values, predictions_test)
r2 = round(r2,5)
print(f"\nR2 score on the test set: {r2}")
```



Final parameter values (theta):

```
[ 1.77406186  4.71824917 18.30394381]
```


First 6 predictions on the test set:

```
[218.10203057 181.49414294 243.42644285 260.3718172  260.07070814
 174.64439034]
```

R2 score on the test set: 0.77806

Actual vs Predict CO2 emissions:

```
# Plot actual vs predicted values
plt.figure(figsize=(8, 6))
plt.scatter(emission_test_values, predictions_test, color='blue', alpha=0.6, label="Predicted vs Actual")
plt.plot([emission_test_values.min(), emission_test_values.max()],
         [emission_test_values.min(), emission_test_values.max()],
         color='red', linestyle='--', label="Ideal Fit")
plt.title("Actual vs Predicted CO2 Emissions")
plt.xlabel("Actual CO2 Emissions")
plt.ylabel("Predicted CO2 Emissions")
plt.legend()
plt.grid(True)
plt.show()
```



## Logistic Regression :Accuracy

```
sgd_classifier = SGDClassifier()
sgd_classifier.fit(independent_features_train, emission_train_classes)

emission_test_classes_pred = sgd_classifier.predict(independent_features_test)

accuracy = accuracy_score(emission_test_classes, emission_test_classes_pred)
print("Accuracy:", accuracy)

Accuracy: 0.9611913357400722
```

## Confusion Matrix:

```
conf_matrix = confusion_matrix(emission_test_classes, emission_test_classes_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=encoder.classes_, yticklabels=encoder.classes_)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

